
Kafka Day 7 – Revision Notes (Exactly-Once Semantics & Reliability)

1. Kafka Partition Ordering Guarantee

- Kafka preserves order **within a partition**.
- If a producer fails mid-transaction, Kafka **aborts** the transaction.
- All writes in that transaction are discarded (atomicity guaranteed).

2. Backpressure Handling in Consumers

- Use `max.poll.records` to limit records fetched per poll.
- Slow downstream? Buffer internally or use async handoff.
- Monitor lag & alert early to prevent snowballing delays.

3. Exactly-Once Processing Semantics

- Use **idempotent producers + transactional producers + atomic offset commit**.
- Commit offsets only **after** DB/sink confirms success.
- Guarantees messages aren't duplicated or lost.

4. Consumer Lag Monitoring

- Lag = `latest offset - committed offset`.
 - Tools: **Burrow, Prometheus, Grafana, Kafka Manager**.
 - Watch out for spikes during rebalancing or processing delays.
-

5. Ensuring Durability and Availability

- Use:
 - `replication.factor = 3`
 - `min.insync.replicas = 2`
 - `acks = all`
 - Disable unclean leader election.
 - Use rack-awareness for better fault tolerance.
-

6. Schema Evolution in Kafka

- Use **Avro/Protobuf + Schema Registry**.
 - Compatibility types:
 - BACKWARD (old consumers work with new schema)
 - FORWARD (new consumers work with old schema)
 - FULL (both ways).
 - Always handle defaults for missing fields.
-

7. Real-Time + Batch Design with Kafka

- Use **separate consumers** for real-time (e.g. Spark Streaming) and batch (e.g. Flink windowed jobs).
 - Use `log.retention.hours` or compacted topics to allow late replays.
-

8. Kafka Monitoring & Alerting

- Monitor:

- Under-replicated partitions
 - Consumer lag
 - Broker disk usage
 - Request throughput
 - Use: **JMX, Prometheus, Grafana, Datadog, Burrow.**
-

9. Ordering Across Partitions

- Kafka **does not guarantee** ordering across partitions.
 - You must design using:
 - Single partition per key
 - Sink-side ordering with buffering (careful of memory pressure).
-

10. Kafka Cluster Migration with Minimal Downtime

- Use **MirrorMaker2** to replicate topics across clusters.
- Validate data parity via topic offsets or hash checksums.
- Gradually shift consumers and producers.