

# Kafka Interview - Day 4 Revision Notes (May 1)

## 1. What is CooperativeStickyAssignor and how does it help?

It enables incremental rebalancing by only revoking partitions that need to be reassigned. This reduces downtime and consumer lag during rebalancing.

## 2. What is the trade-off of increasing session.timeout.ms?

Higher session timeout reduces unnecessary rebalances but delays detection of dead consumers. It's a balance between stability and fault detection.

## 3. How does compression affect producer/consumer performance?

Compression reduces payload size and improves throughput, but adds CPU overhead. Choose codecs based on workload: Snappy for speed, GZIP for size.

## 4. How is exactly-once semantics achieved in Kafka?

By enabling idempotent producers, transactions, and committing consumer offsets within the same transaction boundary using Kafka's transactional APIs.

## 5. How do you monitor and troubleshoot Kafka consumer lag?

Use tools like Burrow or Prometheus. Track metrics like ``consumer_lag``, ``records-lag-max``, and alert when thresholds are breached.

## 6. How to avoid partition skew or uneven data distribution?

Choose the right assignor (e.g., StickyAssignor), use good partitioning keys, and monitor throughput per partition to avoid hotspots.

## 7. What is ISR and how does it ensure durability?

ISR (In-Sync Replicas) are brokers that are fully caught up with the leader. Kafka only acknowledges writes when `min.insync.replicas` are met to ensure durability.

### **8. How do you ensure consistency with multiple consumer groups?**

Use committed offsets, exactly-once delivery (if needed), and ensure schema validation to maintain consistency across multiple consumers.

### **9. How do you achieve fault tolerance in Kafka architecture?**

Use replication, enable retries with backoff, monitor ISR, and have high availability consumers and producers to handle broker or consumer failure.

### **10. How to handle out-of-order messages in Kafka?**

Use partitioning keys to enforce ordering within partitions. For full ordering, use external sequencing or deduplication logic downstream.

### **11. How to reduce end-to-end latency in Kafka pipelines?**

Use compression, batch records, reduce `linger.ms`, avoid excessive rebalancing, and fine-tune producer/consumer fetch configurations.

### **12. How do you monitor Kafka broker performance?**

Track metrics like ``under_replicated_partitions``, ``active_controller_count``, CPU usage, GC time, and use Grafana or Confluent Control Center.